Candidate Signal Score: a BrightHire-powered feedback and fairness layer

Patrick Diamitani - prepared for the Gong Talent Engineer conversation with Robert Liedtka, October 26, 2026

The idea in one line: apply lead-scoring rigor to interview evidence. A transparent score tells the hiring team how well-evidenced each candidate's loop is, flags unfair gaps before they become rejections, and turns interview signals into feedback every candidate actually receives.

Why this build
Hiring funnels fail the same way revenue funnels do: leaky stages, inconsistent data, decisions made on partial signal. Gong's stack already captures the hard part - BrightHire records and structures every interview (AI notes, topic coverage, one-click scorecards, interview-quality metrics) and syncs it into Greenhouse. What's missing is the layer that turns those signals into better decisions and better communication. This workflow is that layer.

How it works
1. Capture. BrightHire interview events (completed interview, submitted scorecard) and Greenhouse stage changes flow into an automation layer (n8n) via webhooks and the Greenhouse Harvest API. BrightHire already writes scorecards, AI notes, and highlight links onto the candidate record, so interviewers do zero new data entry.
2. Store. Every event lands in a governed store - one table each for candidates, stages, interviews, and scorecards - with role-based access from day one. Same pattern as the talent-intelligence repository I'd build across Workday, Greenhouse, Gem, and BrightHire.
3. Score. A transparent scoring function computes each candidate's Signal Score from four auditable inputs:
   - Rubric coverage: share of required competencies with captured evidence
   - Evidence depth: AI-note citations and scorecard substance per competency
   - Consistency: whether candidates for the same role were asked the same core questions
   - Scorecard timeliness: interviewer feedback submitted against SLA
4. Act. Three outputs:
   - Debrief packet: a per-candidate evidence matrix assembled before the debrief, so the panel argues from the same facts.
   - Thin-signal flag: if a candidate is about to be rejected with under 60% rubric coverage, the workflow pauses the rejection and asks for a re-review or a targeted follow-up. Nobody gets rejected because of a gap in our process.
   - Feedback loop: every decision triggers an AI-drafted, human-approved feedback message to the candidate within a 48-hour SLA, built from the actual interview evidence.

Guardrails (the point, not a footnote)
The score never auto-dispositions anyone. It ranks evidence quality and prioritizes human attention; humans make every decision. No demographic or proxy features, ever. Every score shows its inputs, so a candidate, an auditor, or a hiring manager can see exactly why it says

what it says. Candidates are told what is captured and how it is used - the same consent-first model BrightHire already runs.

Why it's measurable
Feedback coverage (% of rejected candidates receiving substantive feedback), time-to-feedback vs SLA, scorecard completion rate, rubric coverage per loop, candidate NPS, and drop-off rate. BrightHire's own study across 25,000+ candidates ties better interview structure to 19% fewer candidate drop-offs and 27% fewer interviews per hire. This layer makes those gains visible and compounding.

Prototype path
Buildable today on free tools: Greenhouse webhook to n8n to Supabase to Slack/email, with recorded-interview transcripts standing in for BrightHire signals in a sandbox. Happy to walk through a working version live.